

OPERATOROS

One platform.
Connected work.
Retained intelligence.

An engineering overview of the platform, its applications, and the workflows that turn operational activity into reusable assets.

13 APPLICATIONS

42 ILLUSTRATIVE SCENARIOS

Prepared for engineering colleagues

John Travis Williams Jr. / Shotgun Ninjas Productions, LLC

SOURCE SNAPSHOT ee9ca05e / 2026-09-29



THE ENGINEERING PROPOSITION

Reduce the cost of context.

OperatorOS centralizes authority while purpose-built applications retain their own work models. The intended value is less re-entry, clearer ownership, and evidence that survives the handoff.

01

Shared authority

Identity, tenants, roles, subscriptions, entitlements, launch policy, and audit have one parent authority.

02

Distinct work models

Service revenue, IT resolution, and healthcare operations remain separate main-module workflows.

03

Reusable outcomes

Reviewed proof, scripts, knowledge, training cases, and campaign packages become durable work products.

READ THIS AS A TECHNICAL BRIEF

Architecture and platform: pages 3–7. **Module inventories and case studies:** pages 8–33. **Resolution Intelligence:** page 34. **Evaluation and evidence:** pages 35–38.

Case-study method: all 42 scenarios are illustrative, not customer testimonials. Expected outputs are workflow goals; measurement suggestions are not reported results. Feature descriptions are grounded in repository contracts and implementation records.

ONE PARENT / THREE MAINS / TEN COMPANIONS

A portfolio with explicit boundaries.

The map shows product scope and access class. Actual access is a server decision; an application category is not a promise that every tenant receives it.

MAIN MODULES

TradeFlowKit
Service work + commercial records

TechDeck
IT context + resolution knowledge

PulseDesk
Healthcare operations coordination

PAID COMPANIONS

BrandForge OS
Brand context + campaign production

SnapProofOS
Field evidence + approved reports

StudyForge AI
Source material + study practice

Deploy Ops
Campaign packages + launch readiness

CallCommand AI
Receptionist + intake + follow-up

Script Ops
Reviewed, versioned script assets

FREE + GATED COMPANIONS

TorqueShed
Vehicle history + diagnostic work

FaultlineLab
Diagnostic practice + investigation traces


Operator Pool Hall
Browser 8-ball + shared recreation

OutCall
Coming soon / production unavailable

Current commerce separates the selected core product, companion grants, and seats. Legacy plan compatibility is a separate path. No pricing quote or all-modules entitlement is implied by this overview.

One authority. Multiple exact hosts.

Next.js web surfaces and a private Fastify API operate as one platform deployment. Module hosts present distinct products; they do not become independent identity or billing systems.



SSO: a 60-second opaque code is single-use and bound to exact callback, state, nonce, PKCE S256, environment, tenant, module, and entitlement. Redemption revalidates authority.

Isolation: host-only Secure/HttpOnly cookies; no ambient parent-domain cookie. Server-resolved tenant predicates govern reads, writes, constraints, transactions, and audit. A requested tenant ID is never authority.

THE REUSABLE FOUNDATION

Capabilities every module should not rebuild.

A common control plane reduces duplicated security and business logic while keeping module records, roles, and outcomes explicit.

Identity and account security

Central sign-in, session lifecycle, role checks, account status, MFA/revocation controls, and exact-host SSO.

Tenants and entitlements

Validated memberships, organization selection, module grants, subscription/seat policy, and server-owned launch decisions.

Shared customer directory

Organization/contact identity reused by supported applications; shared contact edits avoid automatic rewriting of historical reports.

Files, jobs, and exports

Private attachments, configured scanning, durable work queues, usage accounting, export artifacts, and explicit failure/retry states.

Billing and service connections

Platform Stripe authority stays separate from module business payments; stored configuration never proves external delivery.

Audit and collaboration

Platform audit, shared activity/notifications, tenant messenger, administrative visibility, and scoped record-to-record handoffs.

Operational intelligence today: deterministic workday briefs rank recorded exceptions and point to the next action. They do not silently invoice, charge, send, assign, or execute.

FROM ADJACENT APPS TO CONNECTED WORK

Handoffs preserve the original context.

Ten registered workflow contracts connect selected records. Each result remains a specific draft, ticket, attachment, or package with a reviewed source version.

- 01

TradeFlowKit job

→ SnapProofOS proof work
- 02

Approved SnapProofOS PDF

→ Originating TradeFlowKit job
- 03

CallCommand analyzed real call

→ TradeFlowKit lead or customer/job
- 04

CallCommand analyzed real call

→ PulseDesk operations request
- 05

CallCommand analyzed real call

→ TechDeck support ticket

- 06

Resolved TechDeck / PulseDesk work

→ Unpublished FaultlineLab draft
- 07

TorqueShed diagnostic

→ SnapProofOS diagnostic proof
- 08

Verified / resolved TorqueShed work

→ Unpublished FaultlineLab draft
- 09

BrandForgeOS campaign

→ Deploy Ops launch package
- 10

Approved Script Ops revision

→ Non-executing TechDeck runbook

Delivery contract: recheck both applications and source access at queue and delivery; reject changed sources; deduplicate the business operation; retain signed event context and scoped results. Training handoffs require manager review. Live-call handoffs reject simulations.

THREE ILLUSTRATIVE PLATFORM CASE STUDIES

Where shared authority pays off.

These are evaluation scenarios using synthetic or approved data. They do not report customer deployments or measured results.

01

One engineer, several workspaces

SITUATION

An engineer belongs to two organizations and must reach the right module without mixing records or permissions.

WORKFLOW

Authenticate centrally, select the organization, and launch an entitled module. Follow a deep link and then return to My Apps; switching organization requires a fresh module authorization.

EXPECTED OUTPUT

A host-specific session bound to the intended tenant and module, with server-side checks on every protected action.

MEASURE IN A PILOT

Cross-tenant denial; deep-link continuity; absence of repeat credential entry.

02

A customer shared across workflows

SITUATION

A service customer appears in commercial records, a campaign brand, and field proof, each with different work history.

WORKFLOW

Use the supported shared customer selector in TradeFlowKit, BrandForge OS, and SnapProofOS. Update the shared contact intentionally and verify current projections without rewriting approved historical reports.

EXPECTED OUTPUT

One reusable customer identity with independent module records and preserved historical evidence.

MEASURE IN A PILOT

Duplicate identity creation; contact consistency; historical report stability.

03

A retry without a duplicate result

SITUATION

A reviewed cross-module transfer is interrupted after submission and another authorized person attempts the same outcome.

WORKFLOW

Queue the supported handoff with its source version. Inspect status and the destination link, then exercise the allowed retry path; changed sources and revoked access should fail closed.

EXPECTED OUTPUT

A traceable, deduplicated destination result for supported tenant-owned workflows; actor-scoped campaign generation remains distinct.

MEASURE IN A PILOT

Duplicate records; stale-source rejection; recoverable retry behavior.

TradeFlowKit

Keep service work connected to cash.

A shared operations workspace for service businesses: customer context, execution, commercial documents, and collections remain connected throughout the job lifecycle.

01 **Lead and customer management**

Capture leads; qualify and convert opportunities; use shared customer identity, contacts, search, and bounded CSV imports.

03 **Tasks and workflow stages**

Track job-scoped tasks, dependencies, comments, tags, workflow transitions, and personal or shared saved views.

05 **Invoices and payments**

Maintain invoice items, partial/full payment records, balances, numbering, and history; connected settlement needs provider acceptance.

07 **Reporting and accounting handoff**

Inspect workload and revenue from persisted records; export invoice CSV, QuickBooks IIF, and Xero-oriented CSV files.

02 **Jobs and recurring work**

Schedule and assign jobs; retain owners, dates, status, recurring templates, pause controls, and run history.

04 **Quotes and conversion**

Build multi-line quotes; record decisions; convert to jobs or invoices with duplicate-safe server operations.

06 **Customer communication**

Use scoped public views, customer portals, and consent-aware email/SMS outbox actions with explicit delivery state.

08 **Proof and recovery**

Create SnapProofOS work from a job, return the approved PDF, and use guarded archive/restore workflows.

TF

THE WORKFLOW

- 1 **Lead + customer**
- 2 **Job + quote**
- 3 **Work + evidence**
- 4 **Invoice + payment**

RETAINED OUTPUT

One traceable chain from customer request through work, proof, invoice, and recorded settlement.

BOUNDARY Live payments and outbound delivery require configured, accepted providers. Accounting exports are files, not live QuickBooks/Xero synchronization. Recording a sent state is not email delivery.

THREE ILLUSTRATIVE CASE STUDIES

TradeFlowKit in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

The unbilled service call

SITUATION

A field-service team finishes a repair, but the invoice handoff is lost between the technician and office.

WORKFLOW

Link the customer, job, task completion, and quote. Use the workday brief to find completed jobs without an invoice, then review and create the commercial record.

EXPECTED OUTPUT

A job-linked invoice with an owner and an explicit balance; no automatic charge is implied.

MEASURE IN A PILOT

Completed jobs without invoices; days from completion to invoice.

02

Recurring maintenance without duplicate jobs

SITUATION

A contractor repeats the same service at several customer sites and recreates the work manually each month.

WORKFLOW

Define recurring job templates, inspect the next run, and let the supported scheduler create the dated job. Pause a schedule during an outage and inspect its history before resuming.

EXPECTED OUTPUT

Repeatable work records whose schedule/run identity prevents replay from creating duplicates.

MEASURE IN A PILOT

Manual setup time; duplicate jobs per scheduled run.

03

A defensible proof-to-payment handoff

SITUATION

An office needs to show what was done before following up on an unpaid installation.

WORKFLOW

Create SnapProofOS work from the originating job. Review the captured evidence and approved report, return its PDF to that exact job, then review invoice and recorded payment state.

EXPECTED OUTPUT

Commercial context and approved proof remain linked. A report attachment does not assert customer receipt.

MEASURE IN A PILOT

Jobs with linked approved proof; time spent reconstructing an invoice history.

TechDeck

Make resolution knowledge reusable.

An MSP and IT operations console that binds support work to client, configuration, network, procedure, and evidence context. Resolution Intelligence preserves what actually happened, including failed attempts.

01 **Service operations**

Manage client/site-linked tickets, assignments, comments, time, requester updates, and accountable closeout.

03 **Network and IPAM records**

Maintain network, subnet, address, and topology relationships as documented operational context.

05 **Resolution evidence intake**

Import structured closeouts with screened source material, immutable revisions, validation, warnings, failures, and side effects.

07 **Evidence-derived knowledge**

Preview deterministic KB/runbook drafts, retain source citations, link multiple incidents, and block stale-source approval.

02 **Configuration and lifecycle inventory**

Track typed configurations, asset context, warranty/renewal dates, lifecycle changes, and related records.

04 **Versioned documentation**

Edit knowledge articles and runbooks with revisions, backlinks, review, approval, publication, and private attachments.

06 **Exact and full-text retrieval**

Search identifiers and narrative; inspect evidence, components, action chronology, revision history, and native links.

08 **Reports and controlled connections**

Use persisted dashboards, client reports, portal/status workflows, scoped API tokens, and signed outbound webhooks.

TD

THE WORKFLOW

- 1 **Request + system**
- 2 **Actions + evidence**
- 3 **Verify + closeout**
- 4 **Retrieve + reuse**

RETAINED OUTPUT

A resolution record that explains the fix, failed paths, uncertainty, verification, and reusable guidance.

BOUNDARY

Runbooks do not execute scripts. Inventory is not live discovery or endpoint control. Phase 5 semantic retrieval is source-implemented and provider-disabled; ordinary exact/text search is independent.

THREE ILLUSTRATIVE CASE STUDIES

TechDeck in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

A recurring application failure

SITUATION

A second engineer inherits a symptom that resembles an earlier incident, but a copied final fix omits a harmful intermediate step.

WORKFLOW

Search the exact error or component. Review the prior chronology, failed action, observed side effect, warning, and validation before choosing a procedure for the new client issue.

EXPECTED OUTPUT

An evidence-informed plan with uncertainty retained. The previous outcome does not prove the same root cause now.

MEASURE IN A PILOT

Time to relevant evidence; repeated failed actions avoided during a reviewed pilot.

02

A client infrastructure handover

SITUATION

An MSP technician leaves and the replacement must reconstruct dependencies from tickets and isolated notes.

WORKFLOW

Connect the client's sites, configurations, IPAM records, lifecycle dates, and versioned documentation. Link current tickets and supporting evidence, then prepare a client-facing operational report.

EXPECTED OUTPUT

A navigable service context with named records and a reusable procedure trail.

MEASURE IN A PILOT

Time to locate the affected system, owner, and current approved procedure.

03

From one repair to a maintained runbook

SITUATION

A difficult incident is resolved, but the team needs repeatable guidance that will not silently drift from its source.

WORKFLOW

Preview an evidence-derived runbook, review the cited sections, save a draft, and submit it through approval. When source evidence changes, inspect the stale warning and prepare current guidance.

EXPECTED OUTPUT

A reviewed document with source-revision links and access inherited from every linked incident.

MEASURE IN A PILOT

Time from closeout to reviewed guidance; stale-source publications prevented.

PulseDesk

Coordinate the work around care.

Healthcare operations coordination for facilities, equipment, supplies, departments, and vendors. The work model captures operational ownership and response history while minimizing patient information.

01 **Facilities and departments**

Maintain directory-linked organizations, facilities, requesters, departments, and operational context.

03 **Queues, teams, and assignment**

Route work to an accountable team or owner; prioritize through queues, filters, saved views, and bulk actions.

05 **Notes and requester updates**

Keep internal coordination separate from requester-facing replies; retain private attachments and work history.

07 **Knowledge and reporting**

Use operational knowledge, persisted dashboards, configuration, preferences, trend reporting, and exports.

02 **Operational request intake**

Create numbered requests with structured classification, equipment-issue prefill, tags, and supporting records.

04 **Targets and escalation**

Track service targets, time, at-risk work, escalation paths, and response history.

06 **Equipment and vendor coordination**


Associate operational assets with facility issues, supply requests, service vendors, and department needs.

08 **Reviewed handoffs**

Accept operations-only CallCommand requests and prepare privacy-reviewed resolved-work drafts for FaultlineLab.

PD

THE WORKFLOW



RETAINED OUTPUT

A resolved operational request with accountable ownership, response timing, and an inspectable work history.

BOUNDARY No patient chart, clinical decision, EHR/PACS integration, or HIPAA certification is claimed. Direct Microsoft 365, Google Workspace, IMAP, and SendGrid mailbox intake is unavailable in this release.

THREE ILLUSTRATIVE CASE STUDIES

PulseDesk in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

An equipment issue between departments

SITUATION

An imaging department reports an unavailable workstation or equipment accessory; responsibility spans facilities and a vendor.

WORKFLOW

Create an operations-only request, associate the facility and asset, assign an owner, set the service target, and retain vendor coordination and requester updates in the work record.

EXPECTED OUTPUT

A visible accountable path to restoring operational availability without storing a patient chart.

MEASURE IN A PILOT

Unassigned age; response time; handoffs without a named owner.

02

A supply delay before a busy shift

SITUATION

A department has an unresolved supply request with a deadline, while the next shift sees only a message thread.

WORKFLOW

Classify the supply need, associate the department and vendor, record the due target, and use the queue to escalate the delay. Close the request with the actual resolution and timing.

EXPECTED OUTPUT

A shared request history that survives shift changes and shows what remains open.

MEASURE IN A PILOT

Overdue operational requests; time spent reconstructing shift context.

03

Learning from repeated facility incidents

SITUATION

Several resolved facilities requests reveal a recurring troubleshooting mistake useful for staff training.

WORKFLOW

Review the operational history and select a suitable resolved request. Managers with access to both applications review privacy and prepare an unpublished FaultlineLab draft for trainer revision.

EXPECTED OUTPUT

A controlled training starting point; the imported first revision cannot publish directly.

MEASURE IN A PILOT

Repeat incident patterns identified; reviewed training cases produced.

TorqueShed

Preserve the reasoning behind the repair.

A garage and automotive diagnostics workspace for vehicle history, structured tests, workshop records, and shareable proof. Torque Assist adds governed guidance when its provider and credit controls are configured.

01 **Garage and service history**

Record vehicles, mileage, owner concerns, repairs, maintenance, and service chronology.

03 **Torque Assist**

Request context-aware diagnostic guidance through OperatorOS-owned AI and usage/token accounting.

05 **Repair reports and collaboration**

Create workshop exports and repair records; use supported live-bay and sharing workflows.

07 **Marketplace listings**

Search and save listings; contact sellers; retain conversations, expiry, reports, and listing history.

02 **Structured diagnostics**

Organize symptoms, codes, observations, likely causes, tests, findings, and verification.

04 **Build journals and cost context**

Track build stages, parts, labor, costs, notes, and supporting media.

06 **Community**

Maintain profiles, preferences, posts, comments, reactions, follows/blocks, private media, and moderation records.

08 **Proof and training handoffs**

Prepare diagnostic proof in SnapProofOS and reviewed automotive challenge drafts in FaultlineLab.

TS

THE WORKFLOW

- 1 **Concern + codes**
- 2 **Test + observe**
- 3 **Repair + verify**
- 4 **History + proof**

RETAINED OUTPUT

A repair narrative that separates symptoms, measured observations, work performed, and verified outcome.

BOUNDARY Physical tests and repair judgment remain with the technician. Marketplace listings do not imply escrow, shipping, buyer protection, or reputation guarantees. AI and native-device acceptance are separate gates.

THREE ILLUSTRATIVE CASE STUDIES

TorqueShed in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

Diagnose before replacing parts

SITUATION

A shop receives an intermittent vehicle complaint with several plausible causes and incomplete prior history.

WORKFLOW

Record symptoms, codes, and earlier work, choose the next physical test, and capture its findings. Use Assist only as reviewed guidance when configured; record the actual repair and verification.

EXPECTED OUTPUT

A test-driven diagnostic record that explains why a repair was selected.

MEASURE IN A PILOT

Diagnostic steps with recorded findings; repeat visits for the same complaint.

02

A build that outlives its original owner

SITUATION

A custom vehicle changes hands after months of modifications spread across receipts and photos.

WORKFLOW

Organize stages, parts, labor, costs, media, and service entries in the build journal. Produce a workshop export or controlled repair record for the receiving party.

EXPECTED OUTPUT

A coherent modification and service history with traceable supporting material.

MEASURE IN A PILOT

Missing part/service records; time required to reconstruct a modification.

03

A verified diagnosis becomes training

SITUATION

A mechanic resolves a difficult fault that would make a useful diagnostic exercise for apprentices.

WORKFLOW

Finish the diagnostic record with recorded tests and findings. A manager reviews privacy and transfers a training draft to FaultlineLab, then edits the narrative and evidence before publication.

EXPECTED OUTPUT

An unpublished automotive case that retains diagnostic structure without bypassing trainer review.

MEASURE IN A PILOT

Reviewed cases created from real work; learner investigation quality.

FaultlineLab

Evaluate the path to the answer.

A diagnostic challenge environment for technical practice. Versioned scenarios, persistent investigation traces, and server scoring let a reviewer inspect the reasoning path as well as the final answer.

- 01

Challenge catalog
Discover published challenges through search, filters, sorting, authored packs, and daily selection.
- 03

Evidence-driven investigation
Select diagnostic actions, reveal clues, record findings, use progressive hints, and submit a conclusion.
- 05

Assignments and progress
Assign cases, inspect learning progress and badges, and select further practice from observed results.
- 07

Analytics and exports
Review attempts and outcome history; export material for coaching and team discussion.

- 02

Persistent attempts
Resume standard, daily, preview, assignment, and Chaos-mode attempts across refresh and restart.
- 04

Server scoring
Evaluate submitted work against the versioned scenario and retain append-only attempt evidence.
- 06

Authoring lifecycle
Validate, preview, publish, retire, import, and export controlled challenge versions.
- 08

Operational case reuse
Convert reviewed TechDeck, PulseDesk, and TorqueShed records into private first-draft challenges.

BOUNDARY

Training scores are scenario performance, not accredited certification or production authorization. Imported incidents require privacy review and a revised trainer-approved version before publication.

FL

THE WORKFLOW

1

Scenario + clues

2

Investigate + test

3

Submit + score

4

Review + practice

RETAINED OUTPUT

A replayable investigation record showing which evidence the learner inspected and how the answer was reached.

THREE ILLUSTRATIVE CASE STUDIES

FaultlineLab in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

A repeatable technician onboarding exercise

SITUATION

Two new engineers know the terminology, but the lead needs to see how they investigate an ambiguous failure.

WORKFLOW

Assign the same versioned challenge, compare action sequences and missed evidence, and discuss the submitted cause and response using the retained attempt trace.

EXPECTED OUTPUT

A concrete coaching conversation grounded in observed diagnostic choices.

MEASURE IN A PILOT

Missed critical clues; unnecessary actions; score progression on comparable cases.

02

Turn a resolved incident into an exercise

SITUATION

An internal outage contains a useful lesson but also client names and sensitive operational details.

WORKFLOW

Use the supported resolved-work handoff, review privacy, and revise the private draft. Validate and preview the case before publishing a controlled version for the intended learners.

EXPECTED OUTPUT

Reusable practice material with a deliberate editorial and privacy boundary.

MEASURE IN A PILOT

Time from incident review to approved exercise; privacy findings caught before publication.

03

Diagnose under changing conditions

SITUATION

Experienced staff need practice handling incomplete clues and disruption without changing a production environment.

WORKFLOW

Use a suitable Chaos or daily challenge. Retain the investigation trace, inspect the result, and repeat a targeted exercise where the original reasoning missed a signal.

EXPECTED OUTPUT

A safe training record that exposes decision habits under scenario pressure.

MEASURE IN A PILOT

Recovery of investigation after new evidence; repeat occurrence of the same reasoning error.

Operator Pool Hall

A free shared space for the team.

A browser-based 8-ball companion with real gameplay, persistent match state, and protected online rooms. Its engineering interest lies in deterministic simulation, server validation, and recovery across interruptions.

01 **Free Shoot practice**

Practice shot selection, aim, power, and cue-ball control in the Canvas play surface.

03 **Local hot-seat play**

Share a device for turn-based play with visible match and turn state.

05 **Authoritative match checks**

Use durable snapshots/events and independent server re-simulation of submitted actions.

07 **Touch, English, and audio**

Use touch-friendly controls, spin/English, audio options, and performance settings.

02 **CPU matches**

Play against seeded CPU behavior with the supported rules and deterministic physics.

04 **Authenticated online rooms**

Create or join protected rooms through OperatorOS identity and access controls.

06 **Reconnect and expiry**

Recover supported room state after interruptions and enforce room lifecycle limits.

08 **Profiles and local results**

Keep supported profile/result records and use the installable web-app experience.

PH

THE WORKFLOW

- 1 Practice or room
- 2 Aim + play
- 3 Server check
- 4 Resume + finish

RETAINED OUTPUT

A recoverable match with consistent turn and shot state across the supported play surfaces.

BOUNDARY A community benefit with no wagering, paid competition, prizes, or independently certified rankings. Online two-device and performance acceptance should be demonstrated in the target environment.

THREE ILLUSTRATIVE CASE STUDIES

Operator Pool Hall in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

A five-minute practice break

SITUATION

An engineer wants a short break between long support sessions without setting up another account.

WORKFLOW

Launch the free companion through OperatorOS and enter Free Shoot. Practice a repeatable shot while adjusting aim, power, and English with the supported controls.

EXPECTED OUTPUT

Immediate, self-contained practice inside the same account ecosystem.

MEASURE IN A PILOT

Time from launcher to first playable shot; input responsiveness on the chosen device.

02

A teammate disconnects mid-match

SITUATION

Two coworkers play in a protected online room and one browser loses connectivity.

WORKFLOW

Play through normal server-checked actions, reconnect within the supported room lifetime, and inspect the restored turn and match state from durable snapshots and events.

EXPECTED OUTPUT

A recoverable game state rather than an unverified client-only continuation.

MEASURE IN A PILOT

State agreement after reconnect; abandoned matches caused by recoverable disconnects.

03

A practical state-consistency demo

SITUATION

A team wants a small, understandable example of authoritative state and replay validation.

WORKFLOW

Use an authorized test session to examine normal shot submission, server re-simulation, and refresh recovery. Compare visible state across clients without claiming a formal security assessment.

EXPECTED OUTPUT

An accessible demonstration of client input, authoritative checks, and durable state.

MEASURE IN A PILOT

Consistent ball/turn state across refreshes and the two test clients.

BrandForge OS

Give creative work a durable system.

A brand and campaign workspace for reusable context, reviewed assets, and coordinated delivery packages. The catalog uses the compact name BrandForgeOS; both refer to the same application.

01 **Brand HQ**

Maintain brand kits, positioning, voice, offers, and audience personas; associate approved shared customer identity.

03 **Copy Studio and workflows**

Develop copy variants, calls to action, ad prompts, strategy, and creative directions through guided modes.

05 **Review and collaboration**

Manage draft/review/approval states, comments, assets, activity, and recoverable selected logos.

07 **Recorded performance and reports**

Track entered results, recommendations, and campaign reports; create controlled exports.

02 **Campaign planning**

Create campaign briefs, objectives, channels, budgets, timelines, and related landing content.

04 **Logo concepts**

Preview wordmark, lockup, badge, and monogram compositions; export editable SVG and standard/2x PNG.

06 **Calendar and reusable templates**

Organize delivery work and reuse global/custom templates subject to the applicable entitlements.

08 **Launch-package handoff**

Transfer a campaign into Deploy Ops with copy, visual-production briefs, work items, and an approval starting state.

BF

THE WORKFLOW

- 1 **Brand + audience**
- 2 **Campaign + variants**
- 3 **Review + approve**
- 4 **Export + measure**

RETAINED OUTPUT

A reusable brand context and reviewed campaign package with clear ownership and delivery decisions.

BOUNDARY Direct advertising, social, email, analytics, and CRM delivery are unavailable in this release. Canva/Figma use manual file import. Generated assets require review; a scheduled item is not publication.

THREE ILLUSTRATIVE CASE STUDIES

BrandForge OS in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

One offer, several inconsistent messages

SITUATION

An MSP launches a service package, but website copy, sales emails, and social drafts describe different benefits.

WORKFLOW

Define the audience, offer, proof points, and voice once. Build campaign variants from that context, review the claims together, and export the approved materials for each external channel.

EXPECTED OUTPUT

A coherent campaign kit whose variants can be traced back to agreed positioning.

MEASURE IN A PILOT

Review rounds caused by inconsistent claims; time to produce an approved variant.

02

Agency work without cross-client confusion

SITUATION

A small team manages several customers and repeatedly copies brand details between disconnected documents.

WORKFLOW

Associate each brand with the correct shared customer, maintain distinct personas and offers, and retain campaign assets, comments, and review states under the authorized organization.

EXPECTED OUTPUT

Customer-linked creative context and an inspectable approval history.

MEASURE IN A PILOT

Misattributed assets; time required to locate the latest approved brand direction.

03

From campaign approval to a launch package

SITUATION

A campaign is approved, but the coordinator still needs channel copy, production briefs, and assigned launch work.

WORKFLOW

Review the current campaign and queue the Deploy Ops handoff. Inspect the returned package, complete missing production items, and export only after the team checks the resulting materials.

EXPECTED OUTPUT

A coordinated execution package; no advertising spend or external publication occurs automatically.

MEASURE IN A PILOT

Time from approval to handoff; missing deliverables discovered before launch.

SnapProofOS

Turn completed work into reviewable proof.

A field-work and evidence workflow for customer jobs, dated captures, findings, review, and branded reports. Controlled sharing is tied to approved report snapshots rather than unrestricted raw evidence.

01 **Customers, jobs, and assignments**

Select shared customer identity, create proof work, assign field tasks, and retain job context.

03 **File integrity and scanning**

Apply type/signature checks, private storage, scan state, and SHA-256 verification through shared services.

05 **Branded report output**

Generate approved PDF and DOCX documents from the selected report snapshot and template.

07 **Custody and retention**

Retain hash-linked custody history, immutable report context, archive controls, retention, and legal-hold state.

02 **Field evidence capture**

Collect photos, private notes, files, observations, parts, labor, and dated work context.

04 **Findings and review**

Submit evidence, inspect findings, record comments, and move reports through authorized review and approval.

06 **Controlled report shares**

Create expiring and revocable report links; inspect access history without exposing the private source library.

08 **Connected proof packages**

Receive TradeFlowKit job or TorqueShed diagnostic context; return the exact approved PDF to its originating job.

SP

THE WORKFLOW

- 1 **Assign + capture**
- 2 **Organize + review**
- 3 **Approve + generate**
- 4 **Share + retain**

RETAINED OUTPUT

A branded, approved proof package whose included evidence can be traced to the field assignment.

BOUNDARY Storage, scanning, and document workers need acceptance. A share link is not proof of delivery or receipt. Hash-linked history supports integrity review; it is not a legal-admissibility guarantee.

THREE ILLUSTRATIVE CASE STUDIES

SnapProofOS in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

A field installation with incomplete closeout

SITUATION

An engineer completes a network cabinet installation, but photos and labor notes remain scattered across devices.

WORKFLOW

Create the customer job, assign capture work, add dated photos and notes, record findings and costs, and assemble the customer-facing template for review.

EXPECTED OUTPUT

One approved report with supporting job context instead of a loose image folder.

MEASURE IN A PILOT

Jobs missing required evidence; time from site departure to approved report.

02

Review a disputed service outcome

SITUATION

A customer questions what was included in a repair or inspection performed several weeks earlier.

WORKFLOW

Open the authorized job and inspect its captured evidence, comments, custody history, and approved report snapshot. Generate or retrieve the exact document relevant to the reviewed work.

EXPECTED OUTPUT

An inspectable record of what was captured and approved, without assuming the record proves causation.

MEASURE IN A PILOT

Time to reconstruct the work package; gaps in the evidence-to-report chain.

03

Share the report, control the access

SITUATION

A project manager needs to give a customer a finished report while keeping internal notes and raw captures private.

WORKFLOW

Approve the snapshot, generate the requested format, and deliberately create an expiring report share. Deliver it through the team's approved channel and revoke access when appropriate.

EXPECTED OUTPUT

A bounded customer report link with expiry, revocation, and access history.

MEASURE IN A PILOT

Unrestricted evidence links created; share expiry/revocation behavior in the pilot.

StudyForge AI

Make training material easier to use.

A source-based study workspace that turns supplied notes into complete study packs and persistent practice. Built-in creation remains useful without the optional shared AI service.

01 **Folders and study sets**

Organize source material, course/exam context, titles, descriptions, and reusable study sets.

03 **Built-in and optional AI paths**

Use deterministic creation or configured shared AI refinement with validation, provenance, and fallback behavior.

05 **Quizzes and review**

Use server-scored attempts, inspect explanations and missed answers, and retain retry/history records.

07 **Content lifecycle**

Edit, search, filter, duplicate, regenerate, archive, restore, and remove supported study material.

02 **Complete study packs**

Create summaries, key terms, flashcards, multiple-choice questions, short answers, review sheets, and plans.

04 **Flashcard learning**

Practice with keyboard/touch controls and retain known/learning state and session progress.

06 **Plans and progress**

Track date-based study plans, exam countdowns, streaks, usage, and learning trends.

08 **Portable output**

Export study sets and supported progress/results as JSON or entitlement-gated CSV.

SF

THE WORKFLOW

- 1 Supply material
- 2 Build + review pack
- 3 Practice + quiz
- 4 Inspect + repeat

RETAINED OUTPUT

A reusable study pack plus a history of practice and quiz results that identifies material needing another pass.

BOUNDARY Generated material needs source review. AI requires shared-service setup. No outside citations, accredited qualification, automatic mastery judgment, or invented classroom/group workflow is claimed.

THREE ILLUSTRATIVE CASE STUDIES

StudyForge AI in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

Prepare for a platform handover

SITUATION

An engineer must learn an internal service before joining the support rotation.

WORKFLOW

Supply approved, non-secret operating notes, build a study pack, correct the generated material against the source, and use flashcards and a quiz to reveal weak areas.

EXPECTED OUTPUT

A repeatable learning aid with an inspectable source basis and saved practice history.

MEASURE IN A PILOT

Questions missed by topic; preparation time using the same approved material.

02

Turn a long workshop into practice

SITUATION

A technical workshop produces useful notes, but participants rarely revisit the full document.

WORKFLOW

Create a focused set from the approved workshop material. Review terms, short answers, and the review sheet; schedule study around the intended date and inspect the quiz explanations.

EXPECTED OUTPUT

Multiple ways to rehearse the same material without claiming the generated pack is authoritative.

MEASURE IN A PILOT

Study sessions completed; repeated mistakes after source corrections.

03

Keep learning when AI is unavailable

SITUATION

A team wants a study workflow before external AI service approval is complete.

WORKFLOW

Use built-in pack creation, manually refine the material, and practice through the existing flashcard and quiz paths. Export the set for review and retain usage/history in the workspace.

EXPECTED OUTPUT

A functioning study process that does not depend on live model access.

MEASURE IN A PILOT

Completion of the study workflow with AI disabled; content corrections found in review.

Deploy Ops

Make campaign launch work inspectable.

A campaign-package and readiness workspace. Despite the name, its current product scope is coordinated campaign preparation and human-verified delivery records, not software deployment automation.

01 **Business templates and briefs**

Start from reusable templates; define audience, offer, action, tone, channels, and launch date.

03 **Visual-production briefs**

Create up to nine visual briefs, with included plan items made explicit.

05 **Launch execution workspace**

Assign phases, milestones, owners, required files, tasks, and dependencies.

07 **Verified exports**

Download supported text, Markdown, JSON, and CSV packages with checksum and audit context.

02 **Multi-channel campaign packages**

Prepare landing, ad, email, SMS, social, FAQ, flyer, and call-to-action copy.

04 **Versioned artifacts**

Retain saved packages, revisions, source context, approved inputs, and related private assets.

06 **Computed readiness**

Calculate readiness from recorded conditions and approvals; inspect missing prerequisites before a launch transition.

08 **Reviewed completion**

Receive BrandForgeOS campaign handoffs and record externally confirmed completion with a reference.

DO

THE WORKFLOW

- 1 Brief + template
- 2 Package + review
- 3 Tasks + readiness
- 4 Export + confirm

RETAINED OUTPUT

A reviewed campaign kit, explicit launch checklist, and verifiable export for the people using publishing tools.

BOUNDARY Does not deploy software, change DNS, publish sites/ads, send campaigns, or buy media. A readiness score reflects recorded package conditions; external completion needs a person's verified reference.

THREE ILLUSTRATIVE CASE STUDIES

Deploy Ops in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

Launch an MSP service offer

SITUATION

A new managed service needs consistent landing copy, FAQs, ads, and a practical launch checklist.

WORKFLOW

Define the business brief and channels, generate the package, review every claim, assign production tasks, and export the approved material to the people operating the external tools.

EXPECTED OUTPUT

A coordinated campaign package with clear owners and outstanding work.

MEASURE IN A PILOT

Missing launch assets; time from accepted brief to reviewed export.

02

Find the launch blocker before launch day

SITUATION

Copy is finished, but required artwork and approval are missing from a campaign nearing its deadline.

WORKFLOW

Use the execution workspace to inspect required files, dependencies, owners, and computed readiness. Complete and review the missing work before recording the launch transition.

EXPECTED OUTPUT

An evidence-based readiness decision with the unresolved dependency visible.

MEASURE IN A PILOT

Late blockers found before the deadline; launch transitions with incomplete prerequisites.

03

Hand a campaign to another team

SITUATION

A campaign coordinator needs a clean package for an agency or internal publishing team.

WORKFLOW

Review the current package, create a verified export, and use the team's normal delivery channel. After publication, a person checks the external result and records its reference.

EXPECTED OUTPUT

A traceable artifact handoff and separately recorded external confirmation.

MEASURE IN A PILOT

Rework caused by missing copy or context; exports lacking a reviewed version.

CallCommand AI

Make every call end in an owned next action.

A business receptionist, call-intelligence, and MSP intake workspace with simulations, flow design, follow-up records, and explicit provider readiness. Real calls require accepted telephony, AI, and billing configuration.

- 01

Receptionist configuration
Define business knowledge, hours, greetings, behavior, and approved routing destinations.
- 03

Number and launch controls
Manage number setup, readiness checks, capacity, usage, service health, and deliberate go-live controls.
- 05

Follow-up operations
Review leads, tasks, alerts, and staff actions with clear ownership and retained call context.
- 07

MSP intake foundation
Use approved lines, SupportLinks, Directory-backed organization/contact association, local cases, and operator screen-pop.

- 02

Call-flow design and simulation
Prepare conversation paths, test a no-cost simulated call, and correct routing before live enablement.
- 04

Call records and intelligence
Search history; inspect summaries, timelines, transcripts, consent state, outcomes, and reviewed analysis.
- 06

Core-module handoffs
Convert eligible analyzed provider calls into reviewed TradeFlowKit, TechDeck, or operations-only PulseDesk work.
- 08

Provider and audit controls
Use signed callbacks, durable retry/outbox records, kill switches, evidence history, and explicit test/live distinctions.

BOUNDARY

Live Twilio/OpenAI/Stripe and callback acceptance remain separate. Simulations cannot create production call handoffs. BMS delivery and privileged Datto/Graph/AD actions are not accepted by the MSP foundation.



THE WORKFLOW

- 1

Configure + simulate
- 2

Accept live setup
- 3

Capture + analyze
- 4

Review + assign

RETAINED OUTPUT

A call outcome with retained context and a clearly owned next action, distinguished from any test simulation.

THREE ILLUSTRATIVE CASE STUDIES

CallCommand AI in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

Test after-hours intake before activation

SITUATION

An MSP wants consistent after-hours intake but needs to inspect what the receptionist will ask and where it will route work.

WORKFLOW

Configure hours, approved knowledge, and escalation paths. Run simulations, inspect captured answers and routing, and correct the flow before the separate provider go-live checks.

EXPECTED OUTPUT

A reviewed receptionist configuration and test outcome without purchasing or placing a live call.

MEASURE IN A PILOT

Required intake fields captured in simulations; incorrect routes found before activation.

02

A real call becomes a support ticket

SITUATION

After provider acceptance, a caller describes an issue that should enter the technical queue with usable context.

WORKFLOW

Review the completed, analyzed provider call and choose the TechDeck handoff. The system reloads the reviewed version and checks access before creating the destination ticket.

EXPECTED OUTPUT

A linked ticket from an eligible real call; simulator records are rejected.

MEASURE IN A PILOT

Calls without an owner; manual re-entry time; duplicate destination tickets.

03

Separate caller context from privileged action

SITUATION

An MSP caller requests urgent account help, but caller recognition must not become authorization for a reset.

WORKFLOW

Use approved intake channels and organization/contact association to capture the issue and open a local case. Review screen-pop and evidence, then use the independently authorized support process.

EXPECTED OUTPUT

Useful intake context without an automatic identity reset or endpoint action.

MEASURE IN A PILOT

Unassociated cases; quality of issue capture; privileged actions requiring explicit separate approval.

Script Ops

Approve the exact automation asset.

A script-authoring and governance workspace for reusable IT automation. Static analysis, immutable versions, explicit review, and audited downloads preserve the boundary between preparing source and executing it.

- 01

Searchable script library
Organize PowerShell, Python, batch, and shell assets with source/provenance context.
- 03

Immutable versions
Pin the exact revision, content identity, and change context used by subsequent review and download.
- 05

Governed review lifecycle
Submit, approve, reject, or retire scripts under tenant-admin controls with recorded decisions.
- 07

Optional AI drafting
Request drafts through shared AI, idempotency, and usage controls; generated source remains unapproved.

- 02

Authoring and import
Create or import source and record the purpose, prerequisites, expected behavior, risk, and rollback.
- 04

Server static analysis
Inspect findings and risk indicators before submitting a script version for approval.
- 06

Approved download controls
Allow only the approved current version to be downloaded; retain integrity and download audit history.
- 08

TechDeck documentation handoff
Create a non-executing draft runbook and protected file-integrity record from an approved script revision.

BOUNDARY

OperatorOS does not execute this source in the browser, API, server, or customer environment. Static analysis is not proof of safety. Execution and endpoint acceptance belong to a separately authorized tool.

SO

THE WORKFLOW

- 1

Draft + requirements
- 2

Analyze + revise
- 3

Review exact version
- 4

Download + document

RETAINED OUTPUT

A specific approved source version with requirements, review findings, rollback context, and an auditable download.

THREE ILLUSTRATIVE CASE STUDIES

Script Ops in practice.

Problem → workflow → expected output. Measurements are suggested pilot checks, not claimed results.

01

Retire the copied desktop script

SITUATION

Several technicians keep slightly different copies of a maintenance script with unclear origin and prerequisites.

WORKFLOW

Import the reviewed source, document requirements and rollback, run static checks, and submit an exact version for approval. Technicians download the approved current revision for their authorized environment.

EXPECTED OUTPUT

One governed automation asset with a visible version and decision history.

MEASURE IN A PILOT

Untracked variants; downloads without an approved current version.

02

Use AI without auto-approving its output

SITUATION

An engineer wants help drafting a repetitive inventory script but must inspect its behavior before distribution.

WORKFLOW

Request a draft through the configured shared AI service, inspect its source, address static findings, and submit the corrected immutable revision to a tenant administrator.

EXPECTED OUTPUT

A reviewable draft that cannot become an approved download merely because a model generated it.

MEASURE IN A PILOT

Findings corrected before approval; generated drafts distributed without review.

03

Preserve procedure beside the code

SITUATION

A useful script exists, but new staff do not know when to use it or how to interpret its output.

WORKFLOW

Select an approved current revision and create the supported TechDeck draft runbook. Review its prerequisites and handling instructions in TechDeck before publishing the procedure.

EXPECTED OUTPUT

Linked source identity and reviewed guidance; the handoff performs no execution.

MEASURE IN A PILOT

Approved assets with current runbook context; source/runbook version mismatches.

COMPANION APPLICATION / COMING SOON / UNAVAILABLE

OutCall

A private, verified-self call workflow.

A reconstructed application for a discreet call to a person's own verified phone. Source/local workflow evidence exists, but the production catalog remains unavailable and these scenarios describe gated behavior.

- 01

Safety acknowledgement
Explain the personal-use and non-emergency boundary before preparing a call workflow.
- 02

Verified-self phone setup
Use a controlled phone-ownership verification path; arbitrary destinations are excluded.
- 03

Neutral call profiles
Store private profiles with controlled voice/script and supported keypad behavior.
- 04

Private exact triggers
Bind exact-match trigger phrases to an approved profile and verified self-owned destination.
- 05

Immediate and scheduled requests
Prepare bounded call requests using durable state and supported scheduling controls.
- 06

Cancellation and history
Inspect actual request/provider state, cancel eligible requests, and retain private history.
- 07

Signed provider events
Enforce controlled voice/SMS/DTMF callbacks, replay checks, and persistent rate limits.
- 08

Private export and deletion
Use password-confirmed export and account-slice deletion controls for the retained personal records.

OC

THE WORKFLOW

- 1

Verify self
- 2

Choose neutral profile
- 3

Request or schedule
- 4

Inspect + cancel

RETAINED OUTPUT

A private request directed only to the verified owner, with honest state and cancellation/history controls.

BOUNDARY

UNAVAILABLE: no live launch or sale. Requires production activation and controlled real-provider acceptance. No emergency response, monitoring, location tracking, duress detection, or third-party dialing.

THREE ILLUSTRATIVE CASE STUDIES

OutCall in practice.

Gated scenarios only. OutCall remains unavailable; these are future acceptance examples, not present live service.

01

A planned social exit cue

SITUATION

A person wants a neutral call to their own phone at a chosen time as a private reason to step away.

WORKFLOW

In a future accepted live workflow, acknowledge the limits, verify the self-owned number, choose a neutral profile, schedule the request, and inspect the actual provider outcome.

EXPECTED OUTPUT

A self-directed call request. This scenario is illustrative; live availability is still blocked.

MEASURE IN A PILOT

Correct verified destination; truthful delivery state; no unauthorized destination expansion.

02

Cancel when the plan changes

SITUATION

A previously scheduled personal call is no longer wanted and should not proceed if it remains cancellable.

WORKFLOW

Open the private history, inspect the request state, and cancel the eligible schedule. Verify that the record distinguishes cancellation from an already submitted provider action.

EXPECTED OUTPUT

An explicit lifecycle decision with no fabricated claim that an in-flight call was recalled.

MEASURE IN A PILOT

Cancellation state consistency; late or duplicate requests in controlled acceptance tests.

03

Review and remove personal history

SITUATION

A user wants a private copy of retained records and then wants the supported account slice removed.

WORKFLOW

After activation acceptance, use the password-confirmed export and deletion controls. Review the product's actual retained scope and verify the resulting state in a controlled test.

EXPECTED OUTPUT

A deliberate privacy workflow; no emergency or third-party-monitoring capability is implied.

MEASURE IN A PILOT

Export access controls; deletion-scope accuracy; resistance to cross-account access.

KNOWLEDGE WITH PROVENANCE AND LIMITS

Preserve evidence before adding inference.

TechDeck distinguishes observed source evidence, deterministic projections, human approval, and optional model-based similarity. Those are different kinds of authority.



DEPLOYED EVIDENCE: v64

Repository records publication and owner read checks for import/prompt/search/linked-knowledge surfaces. Production writes, other-role checks, live logout, and restore remain separate acceptance scopes.

SOURCE CANDIDATE: v65 / PROVIDER OFF

OpenAI adapter and bounded hybrid retrieval are implemented. Provider activation, production vector capability, real-model relevance, and publication are not established. Similarity is not repair confidence.

A PROPOSED 45-MINUTE TECHNICAL WALKTHROUGH

Prove the transitions, then discuss scale.

Suggested agenda, not a completed test. Use synthetic records in an approved environment with explicit module grants and a second tenant for denial checks.

00–07	Authority and navigation Launch TechDeck through OperatorOS, follow a deep link, return to My Apps, and inspect a foreign-tenant denial.
07–19	Incident to retained intelligence Import a synthetic closeout, search an exact identifier, review a failed attempt, and preview a source-cited runbook draft.
19–28	An approved asset handoff Inspect an approved Script Ops revision and its non-executing TechDeck draft; show source identity and review status.
28–37	Evidence to work context Demonstrate a TradeFlowKit-to-SnapProofOS proof workflow only in a separately entitled context; inspect the returned approved PDF.
37–45	Failure behavior and discussion Review stale-source rejection, retry identity, unavailable-provider states, deployment evidence, and open acceptance gates.

Discussion prompts: Where does authority reside? Which handoff reduces the most re-entry? Which result is a draft versus an external action? What must a production pilot prove before wider use?

IMPLEMENTATION, DEPLOYMENT, AND ACCEPTANCE

Treat evidence as a ladder.

This document is a repository-grounded product and architecture review, not a new deployment audit. Historical records are attributed to their recorded scope and date.

SOURCE

Observed in this review

Current checkout ee9ca05e contains the catalog, contracts, APIs, UI, and documentation used here.

LOCAL TESTS

Reported in repository

Phase 5 records 14/14 local release stages, 1,593 API, 108 integration, 52 unit, 32 browser, and 4 visual checks; 36 vector-focused checks are separate.

RECORDED LIVE

Specific v64 evidence

The 2026-09-27 record reports v64/64, a published build, 47/47 public checks, and owner read/SSO/deep-link checks. No new live checks were run for this PDF.

PROVIDER / WRITE

Acceptance remains scoped

Live calls, payments, messages, model quality, production writes, other roles, data migration, and restore need their own evidence.

UNAVAILABLE

An explicit product state

OutCall remains coming soon. New semantic retrieval is provider-disabled. A configured key, a green build, or a rendered page cannot substitute for acceptance.

REFERENCE REGISTER 1 / 2

Trace the claims to the repository.

Clickable references are pinned to source commit ee9ca05e. Repository access may be required. Newer release records and current product/registry contracts supersede older status prose.

S01	Canonical catalog and product outcomes Names, hierarchy, access classes, coming-soon state; paired with product-value.ts and products.ts.
S02	Module capability and integration contracts Primary workflow, retained deliverables, provider boundaries, and explicit exclusions for all 13 applications.
S03	Runtime registry and exact-host policy Launch authority and default OutCall denial; test-only launch is explicitly bounded.
S04	OperatorOS SSO contract v1 Opaque codes, 60-second expiry, PKCE S256, host-only sessions, role/tenant binding, revocation.
S05	Ecosystem integration contract Shared authority, tenant enforcement, billing separation, shared-service boundaries.
S06	Current release gate and live evidence Recorded v64 publication/read checks and separate v65 candidate/provider boundary.
S07	Ten connected outcomes Supported handoffs, source freshness, current access, deduplication, and external-action limits.

S08	Current implementation evidence Dated local checks and limitations; reported historical evidence is not a new test run.
S09	Evidence-derived TechDeck documents Deterministic assembly, source citations, review, inherited access, and stale-source protection.
S10	TechDeck semantic-search guide Disabled provider, excerpt review, explicit opt-in, limits, fallback, and no real-model acceptance.
S11	Shared customers and account security Customer links, shared contact edits, account controls, provider truth, and validation scope.
S12	Current parity index Recent overlays and per-surface evidence; newest applicable status takes precedence.
M01	TradeFlowKit capability matrix Jobs, revenue documents, recurring work, exports, and supported recovery.
M02	TechDeck capability matrix Service records, configuration, network, documentation, reporting; supplement with S09/S10.

Trace the claims to the repository.

Clickable references are pinned to source commit ee9ca05e. Repository access may be required. Newer release records and current product/registry contracts supersede older status prose.

M03	PulseDesk capability matrix Healthcare operations scope, queue ownership, equipment/facility/supply workflows.
M04	TorqueShed capability matrix Garage, diagnostics, Assist, community, marketplace, and verification boundaries.
M05	FaultlineLab capability matrix Versioned challenges, attempts, scoring, authoring, assignments, and exports.
M06	Operator Pool Hall capability matrix Supported play modes, durable state, server simulation, and exclusions.
M07	BrandForgeOS capability matrix Campaign and brand lifecycle; current direct-provider exclusions are specified in S02.
M08	SnapProofOS capability matrix Evidence and custody foundation; current reports/shares and connected outcomes use S02/S07.
M09	StudyForge AI capability matrix Complete study packs, persistent learning, authoritative quiz scoring, export, and exclusions.

M10	Deploy Ops capability matrix Campaign artifacts, dependencies, readiness, checksum exports, and publication boundary.
M11	CallCommand AI and MSP intake MSP intake constraints; current receptionist scope and provider setup use S02.
M12	Script Ops capability matrix Versioned scripts, static checks, review, downloads, and non-execution boundary.
M13	OutCall capability matrix Authorized reconstruction is source/local evidence; production remains coming soon.
S13	Runtime topology and deployment model Shared Next.js/Fastify runtime, private API, host routing; registry overrides stale OutCall summary.
S14	Module consolidation record Historical restoration detail; later product contracts and current release evidence control current claims.
S15	Backup and restore contract One ordered database authority, backup/apply/verify flow, and separate restore acceptance.

Editorial precedence: current registry and product contracts resolve stale descriptions (notably OutCall, Deploy Ops, and direct provider connections). Historical module matrices describe implementation provenance; they do not grant new live acceptance.